

SHOULD HAVE A BETTER TITLE.

Anonymous authors

Paper under double-blind review

ABSTRACT

We proposed a novel paradigm for implementing latent reasoning by combining DAPO-like reinforcement learning and special gating units. In this way, we transformed text-based explicit reasoning models to latent reasoning models, progressively and practically. Initial experiments demonstrate its feasibility, with significant stability and the least extra training overhead by parallelism and minimal changes, while holding full potential to scale up.

1 INTRODUCTION

The Chain of Continuous Thought (COCONUT) (Hao et al., 2024) explores the potential of the latent layer of reasoning as an alternative to the traditional explicit CoT (Wei et al., 2022) in both the reasoning procedure and the response procedure. Its previous work has shown that the latent layer of continuous reasoning can reduce unnecessary statements and focus on the core of the problem.

These models are issued from general-purpose language, trained in supervised steps of a reasoning procedure, and progressively increase their layers of latent reasoning. It is impractical due to the lack of customized reasoning data, and progressive training is significantly less convenient than explicit CoT. Hence, we provide a new approach to implementing this latent reasoning layer, fine-tuning from explicit reasoning models.

In 2014 (Chung et al., 2014), Chung et al. proposed the Gated Recurrent Unit. It is first leveraged to improve the RNNs, but we found that using it in the latent reasoning chain is also feasible to maintain its long-term contextual awareness.

In January, the DeepSeek team proposed a distillation-based method to implement a reasoning model through reinforcement learning (DeepSeek-AI, 2025), and applied the approach to Qwen-2.5B (Yang et al., 2024a) to a reasoning model. Given the limitations of a 2.5B parameter model in reasoning tasks, we aim to implement COCONUT on the small model and find its potential in reasoning compared to the original explicit reasoning model.

2 RELATED WORK

Explicit Reasoning Models Advances in explicit reasoning models, including OpenAI o1 & o3 (OpenAI, 2025; 2024) and DeepSeek R1 (DeepSeek-AI, 2025), demonstrates its strong ability on mathematical reasoning. Based on Chain-of-Thought (Wei et al., 2022), just through prompts, models achieved significant improvements on arithmetic, commonsense, and symbolic reasoning tasks with *explicitly generate* intermediate reasoning text. Kojima et al. (2023) discovered that even without few-shot examples, math reasoning capabilities can be still improved by a phrase like “Let’s think step by step.” Yao et al. (2023b) proposed REACT, which combines reasoning and acting, and achieved state-of-the-art performance on several reasoning tasks.

Several novel paradigms were also issued from explicit reasoning models (Wang et al., 2023; Yao et al., 2023a; Besta et al., 2024). In this way, explicit reasoning creates a new level on mathematical reasoning tasks.

Chain of Continuous Thought Previous works already found LLMs may involve hidden computation steps, which are not explicitly generated in text (Yang et al., 2024b; Chen et al., 2023). Driven by this spirit, Hao et al. (2024) proposed COCONUT, assigning models to “reason in a latent continuous space” through specifically designed training objectives, and demonstrated its potential

compared to explicit Chain-of-Thought reasoning procedure, especially on tasks requiring back-tracking. Geiping et al. (2025) introduced a new recurrent transformer architecture that lets models perform multiple hidden reasoning iterations per token. In this way, recurrent block that can be looped arbitrarily at inference to refine the hidden state before producing the next token. Chen et al. (2025) argued that LLMs may involve born ability to reason internally, suggesting LLMs can internally learn better reasoning policies in latent space when given appropriate training signals.

Compared to explicit reasoning, latent reasoning paradigms leverages latent spaces, promoting its efficiency and accuracy, but still face problems with scaling-up and stability.

Large-scale Reinforcement Learning Ouyang et al. (2022) showed how GPT-3 can be fine-tuned with reinforcement learning via Proximal Policy Optimization (PPO) to make it possible to follow human instructions. PPO (Schulman et al., 2017) is a reinforcement learning algorithm that optimizes the policy of an agent by maximizing the expected reward while ensuring that the new policy does not deviate too much from the old policy. Rafailov et al. (2024) proposed a new method called Direct Preference Optimization (DPO), which is a simple and efficient alternative to PPO for fine-tuning language models. Anthropic explored scaling RLHF by leveraging AI models as human labelers. In their Constitution AI Bai et al. (2022), they used RLHF to train a model to significantly increase its safety and harmlessness via reinforcement learning. Shao et al. (2024) proposed Group Relative Policy Optimization (GRPO), which is a new reinforcement learning algorithm that overcame the limitations of PPO, including the need for a large number of samples and the difficulty of tuning hyperparameters. Yu et al. (2025) demonstrated the effectiveness of Decoupled Clip and Dynamic Sampling Policy Optimization (DAPO) in large-scale reinforcement learning, which made state-of-the-art performance on several benchmarks for Qwen2.5-32B Yang et al. (2024a), which scored 50 in AIME 24.

3 METHODOLOGY

3.1 MODEL ARCHITECTURE

Modern LLMs involves Embedding, Transformer, and Output layers, as we denote them E , T_n , and O respectively, where n indicates the layer of hidden state.

To avoid breaking special usage in final layers (Elhage et al., 2021), we used middle layers as the entry of the latent reasoning space.

3.2 INFERENCE

The whole inference procedure is described in (...)

After extracting the hidden state $h_{n,0}$ (where n denotes the layer number and k denotes time sequence, which is 0 here), we pass the state to the Variational Autoencoder (Kingma and Welling, 2019), and produce a vector for storage with length l_s .

Then, the new vector with the same shape of $h_{n,k}$, as we denote $h'_{n,k}$, will be operated by several gates, as described in Equation 1.

$$h_{n,k} = h_{n,k-1} \odot g_k + (1 - g) \odot t_k \quad (1)$$

Gates There are three gates in our model, which are the general gate, the injection scalar, and the Dynamic Current Step Mixing Gate.

1. **General Gate.** The general gate is derived from sequence, its value dynamic. We denote it as g_t , and its procedure of derivation can be described in Equation (not yet).
2. **Injection Scalar.** The injection scalar is a constant value, which brought the result of gate-shaper to the gate, to bridge the difference in order of magnitude.
3. **Dynamic Current Step Mixing Gate.**

Long-term Cache To enable long-term storage of reasoning features, we introduce a VAE-style bottleneck layer, compressing 1536-dimensional features to 256 dimensions and then expanding back to 1536 before gating.

We store it in each operation, and concatenate these vectors into a matrix subsequently as hidden state.

3.3 TRAINING

We use the DAPO objective as our reinforcement learning framework, but the whole procedure is slightly different from the original DAPO.

3.3.1 REINFORCEMENT LEARNING

The whole procedure of reinforcement learning can be described in Figure.

We adapt the DAPO objective, which is formulated as:

$$\begin{aligned} \mathcal{J}_{\text{DAPO}}(\theta) = & \mathbb{E}_{(q,a) \sim \mathcal{D}, \{o_i\}_{i=1}^G \sim \pi_{\theta_{\text{old}}}(\cdot|q)} \\ & \left[\frac{1}{\sum_{i=1}^G |o_i|} \sum_{i=1}^G \sum_{t=1}^{|o_i|} \min \left(r_{i,t}(\theta) \hat{A}_{i,t}, \text{clip}(r_{i,t}(\theta), 1 - \varepsilon_{\text{low}}, 1 + \varepsilon_{\text{high}}) \hat{A}_{i,t} \right) \right] \\ \text{subject to } & 0 < |\{o_i \mid \text{is_equivalent}(a, o_i)\}| < G, \end{aligned} \quad (2)$$

where

$$r_{i,t}(\theta) = \frac{\pi_{\theta}(o_{i,t} \mid q, o_{i,<t})}{\pi_{\theta_{\text{old}}}(o_{i,t} \mid q, o_{i,<t})}, \quad \hat{A}_{i,t} = \frac{R_i - \text{mean}(\{R_i\}_{i=1}^G)}{\text{std}(\{R_i\}_{i=1}^G)}. \quad (3)$$

Firstly we apply rigorous datasets to the inference procedure, and verify the answer (Kydliček).

We parallelized several inference procedures simultaneously, preferably 16 for a same question, and verify the answer.

Award Mechanism We use $\text{award}_{\text{acc}}$ and $\text{penalty}_{\text{acc}}$ to award or penalize in accuracy. After mixing the accuracy and length evaluation, we inject it to loss.

Positive-Negative Ratio We use a ratio to mix positive and negative cases, which is 1 : 1.

3.4 TIME CONSISTENCY REGULARIZATION

A significant reason for chaotic response of multi-layer COCONUT is its time inconsistency, which lead to the model’s failure to maintain a consistent reasoning process.

We adapt a mechanism called “Gating Value Bonus,” which is used to guarantee the gate to converge to a fixed position, as we denote in Equation 4.

$$\mathcal{J}_{\text{gate}} = \min \left\{ C_0 \exp \left[\lambda \left(\frac{1}{2} - g_i \right)^2 \right] \right\} \quad (4)$$

3.5 PARALLELIZED TRAINING

After sampling, we concatenate the question, answer, and hidden state, and apply it to training procedure.

4 EXPERIMENTS

5 DISCUSSION

6 CONCLUSION

REFERENCES

- Y. Bai, S. Kadavath, S. Kundu, A. Askell, J. Kernion, A. Jones, A. Chen, A. Goldie, A. Mirhoseini, C. McKinnon, C. Chen, C. Olsson, C. Olah, D. Hernandez, D. Drain, D. Ganguli, D. Li, E. Tran-Johnson, E. Perez, J. Kerr, J. Mueller, J. Ladish, J. Landau, K. Ndousse, K. Lukosuite, L. Lovitt, M. Sellitto, N. Elhage, N. Schiefer, N. Mercado, N. DasSarma, R. Lasenby, R. Larson, S. Ringer, S. Johnston, S. Kravec, S. E. Showk, S. Fort, T. Lanham, T. Telleen-Lawton, T. Conerly, T. Henighan, T. Hume, S. R. Bowman, Z. Hatfield-Dodds, B. Mann, D. Amodei, N. Joseph, S. McCandlish, T. Brown, and J. Kaplan. Constitutional ai: Harmlessness from ai feedback, 2022. URL <https://arxiv.org/abs/2212.08073>.
- M. Besta, N. Blach, A. Kubicek, R. Gerstenberger, M. Podstawski, L. Gianinazzi, J. Gajda, T. Lehmann, H. Niewiadomski, P. Nyczyk, and T. Hoefler. Graph of thoughts: Solving elaborate problems with large language models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(16):17682–17690, Mar. 2024. ISSN 2159-5399. doi: 10.1609/aaai.v38i16.29720. URL <http://dx.doi.org/10.1609/aaai.v38i16.29720>.
- H. Chen, Y. Feng, Z. Liu, W. Yao, A. Prabhakar, S. Heinecke, R. Ho, P. L. Mui, S. Savarese, C. Xiong, and H. Wang. Language models are hidden reasoners: Unlocking latent reasoning capabilities via self-rewarding, 2025. URL <https://openreview.net/forum?id=4Po8d9GAfQ>.
- W. Chen, M. Yin, M. Ku, P. Lu, Y. Wan, X. Ma, J. Xu, X. Wang, and T. Xia. Theoremqa: A theorem-driven question answering dataset, 2023. URL <https://arxiv.org/abs/2305.12524>.
- J. Chung, Ç. Gülçehre, K. Cho, and Y. Bengio. Empirical evaluation of gated recurrent neural networks on sequence modeling. *CoRR*, abs/1412.3555, 2014. URL <http://arxiv.org/abs/1412.3555>.
- DeepSeek-AI. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning, 2025. URL <https://arxiv.org/abs/2501.12948>.
- N. Elhage, N. Nanda, C. Olsson, T. Henighan, N. Joseph, B. Mann, A. Askell, Y. Bai, A. Chen, T. Conerly, N. DasSarma, D. Drain, D. Ganguli, Z. Hatfield-Dodds, D. Hernandez, A. Jones, J. Kernion, L. Lovitt, K. Ndousse, D. Amodei, T. Brown, J. Clark, J. Kaplan, S. McCandlish, and C. Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. <https://transformer-circuits.pub/2021/framework/index.html>.
- J. Geiping, S. McLeish, N. Jain, J. Kirchenbauer, S. Singh, B. R. Bartoldson, B. Kailkhura, A. Bhatele, and T. Goldstein. Scaling up test-time compute with latent reasoning: A recurrent depth approach, 2025. URL <https://arxiv.org/abs/2502.05171>.
- S. Hao, S. Sukhbaatar, D. Su, X. Li, Z. Hu, J. Weston, and Y. Tian. Training large language models to reason in a continuous latent space, 2024. URL <https://arxiv.org/abs/2412.06769>.
- D. P. Kingma and M. Welling. An introduction to variational autoencoders. *CoRR*, abs/1906.02691, 2019. URL <http://arxiv.org/abs/1906.02691>.
- T. Kojima, S. S. Gu, M. Reid, Y. Matsuo, and Y. Iwasawa. Large language models are zero-shot reasoners, 2023. URL <https://arxiv.org/abs/2205.11916>.
- H. Kydlíček. Math-Verify: Math Verification Library. URL <https://github.com/huggingface/math-verify>.
- OpenAI. Introducing openai o1, September 2024. URL <https://openai.com/o1/>. Accessed: 2025-04-27.

- OpenAI. Introducing openai o3 and o4-mini, April 2025. URL <https://openai.com/index/introducing-o3-and-o4-mini/>. Accessed: 2025-04-27.
- L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. L. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, J. Schulman, J. Hilton, F. Kelton, L. Miller, M. Simens, A. Askell, P. Welinder, P. Christiano, J. Leike, and R. Lowe. Training language models to follow instructions with human feedback, 2022. URL <https://arxiv.org/abs/2203.02155>.
- R. Rafailov, A. Sharma, E. Mitchell, S. Ermon, C. D. Manning, and C. Finn. Direct preference optimization: Your language model is secretly a reward model, 2024. URL <https://arxiv.org/abs/2305.18290>.
- J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal policy optimization algorithms. *CoRR*, abs/1707.06347, 2017. URL <http://arxiv.org/abs/1707.06347>.
- Z. Shao, P. Wang, Q. Zhu, R. Xu, J. Song, X. Bi, H. Zhang, M. Zhang, Y. K. Li, Y. Wu, and D. Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models, 2024. URL <https://arxiv.org/abs/2402.03300>.
- X. Wang, J. Wei, D. Schuurmans, Q. Le, E. Chi, S. Narang, A. Chowdhery, and D. Zhou. Self-consistency improves chain of thought reasoning in language models, 2023. URL <https://arxiv.org/abs/2203.11171>.
- J. Wei, X. Wang, D. Schuurmans, M. Bosma, E. H. Chi, Q. Le, and D. Zhou. Chain of thought prompting elicits reasoning in large language models. *CoRR*, abs/2201.11903, 2022. URL <https://arxiv.org/abs/2201.11903>.
- A. Yang, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Li, D. Liu, F. Huang, H. Wei, H. Lin, J. Yang, J. Tu, J. Zhang, J. Yang, J. Yang, J. Zhou, J. Lin, K. Dang, K. Lu, K. Bao, K. Yang, L. Yu, M. Li, M. Xue, P. Zhang, Q. Zhu, R. Men, R. Lin, T. Li, T. Xia, X. Ren, X. Ren, Y. Fan, Y. Su, Y. Zhang, Y. Wan, Y. Liu, Z. Cui, Z. Zhang, and Z. Qiu. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024a.
- S. Yang, E. Gribovskaya, N. Kassner, M. Geva, and S. Riedel. Do large language models latently perform multi-hop reasoning?, 2024b. URL <https://arxiv.org/abs/2402.16837>.
- S. Yao, D. Yu, J. Zhao, I. Shafran, T. L. Griffiths, Y. Cao, and K. Narasimhan. Tree of thoughts: Deliberate problem solving with large language models, 2023a. URL <https://arxiv.org/abs/2305.10601>.
- S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. React: Synergizing reasoning and acting in language models, 2023b. URL <https://arxiv.org/abs/2210.03629>.
- Q. Yu, Z. Zhang, R. Zhu, Y. Yuan, X. Zuo, Y. Yue, T. Fan, G. Liu, L. Liu, X. Liu, H. Lin, Z. Lin, B. Ma, G. Sheng, Y. Tong, C. Zhang, M. Zhang, W. Zhang, H. Zhu, J. Zhu, J. Chen, J. Chen, C. Wang, H. Yu, W. Dai, Y. Song, X. Wei, H. Zhou, J. Liu, W.-Y. Ma, Y.-Q. Zhang, L. Yan, M. Qiao, Y. Wu, and M. Wang. Dapo: An open-source llm reinforcement learning system at scale, 2025. URL <https://arxiv.org/abs/2503.14476>.

A APPENDIX